

Lossless join



SAPIENZA
UNIVERSITÀ DI ROMA

What does lossless join



usually, a schema is decomposed:

- **when it is not in 3NF**
- to improve efficiency
 - the smaller the size of the tuples, the greater the number of tuples we can load into memory in the same read operation
 - the tuple information is not used entirely by the operations on the database
 - example: the Student schema could be decomposed by separating the personal information (TaxC, Name, Surname, DateB, PlaceB, etc.) from the academic information (Matriculation, Degree, Year, etc.)



- we have seen that when a schema is decomposed, **it is not enough for the** subschemas to be in 3NF
- let's see two examples (one more "abstract", the other more "concrete")

What we want to achieve - 3NF is not enough



- we now consider the schema $R=ABC$ with the set of functional dependencies $F=\{A \rightarrow B, C \rightarrow B\}$ (the schema is not in 3NF due to the presence in F^+ of the partial dependencies $A \rightarrow B$ and $C \rightarrow B$, since the key is AC)
- this schema can be decomposed into:
 - $R1=AB$ with $A \rightarrow B$
 - $R2=BC$ with $C \rightarrow B$
- the resulting schemas, even if they **preserve all dependencies in F^+** , is still not satisfactory.

What we want to achieve - 3NF is not enough



- consider the **legal** instance of R :

R	A	B	C
	a1	b1	c1
	a2	b1	c2

the two facts $(a1, b1, c1)$ and $(a2, b1, c2)$ are true and not others

- we decompose it by obtaining:

R1	A	B
	a1	b1
	a2	b1

R2	B	C
	b1	c1
	b1	c2

What we want to achieve - 3NF is not enough



- if we now compute the natural join we obtain:

R

A	B	C
a1	b1	c1
a2	b1	c2
a1	b1	c2
a2	b1	c1

} these tuples did not exist in the original instance!

- **we must ensure that the decomposition and the following natural join does not result in any loss of information**

What we want to achieve - 3NF is not enough



- consider the schema:
- $R = \{EmployeeID, ProjectID, Manager\}$ with the set of functional dependencies :
- $F = \{EmployeeID \rightarrow ProjectID, ProjectID \rightarrow Manager\}$
- a project can have multiple managers but each manager has only one project, and an employee on a project reports to only one manager
- the schema is not in 3NF due to the presence in F^+ of the partial dependencies $EmployeeID \rightarrow ProjectID$ and $ProjectID \rightarrow Manager$, since the key is $(EmployeeID, Manager)$
- the schema can be decomposed into:
 - $R1 = \{Matriculation, Project\}$ with $EmployeeID \rightarrow ProjectID$ and
 - $R2 = \{Project, Boss\}$ with $ProjectID \rightarrow Manager$
- such a schema, **while preserving all dependencies in F^+** , is not satisfactory

What we want to achieve - 3NF is not enough



- consider the **legal instance** of R :

R	EmployeeID	ProjectID	Manager
	501	30	E1
	502	30	E2

only the two facts (501,30,E1) and (501,30,E2) are true!

- based on the given decomposition, this instance decomposes into:

R1	EmployeeID	ProjectID
	501	30
	502	30

R2	ProjectID	Manager
	30	E1
	30	E2

What we want to achieve - 3NF is not enough



- ...and instead if you join the two legal instances resulting from the decomposition you get

R	EmployeeID	ProjectID	Manager
	501	30	E1
	502	30	E2
	501	30	E2
	502	30	E1



tuples unrelated to the reality of interest

What we want to achieve - 3NF is not enough



in conclusion, the following requirement of the decomposed schema should be met:

- the decomposition must allow us to **reconstruct, by natural join, each legal instance of the original schema** (without adding any extraneous information)

if we decompose a relation schema R we want the obtained **decomposition** $\rho = R_1, R_2, \dots, R_k$ such that each **legal instance** r of R can be **reconstructed through natural join** from the legal instances r_1, r_2, \dots, r_k of the decomposition schemas R_1, R_2, \dots, R_k

as for **reconstructing a tuple t of r** it is required that $t[R_i] \in r_i, \forall i = 1, \dots, k$, then we must have $\pi_{R_i}(r) = r_i, \forall i = 1, \dots, k$

Definition

let R be a relation schema; A decomposition $\rho = R_1, R_2, \dots, R_k$ has a **lossless join** if for each legal instance r of R we have

$$r = \pi_{R_1}(r) \bowtie \pi_{R_2}(r) \bowtie \dots \bowtie \pi_{R_k}(r)$$



- again, we start with a given decomposition, and look for a way to verify that it complains with the definition we gave

Theorem



Theorem: let R be a relation scheme and let $\rho = R_1, R_2, \dots, R_k$ be a decomposition of R ; for each legal instance r of R , denoted by $m_\rho(r) = \pi_{R_1}(r) \bowtie \pi_{R_2}(r) \bowtie \dots \bowtie \pi_{R_k}(r)$, we have:

- $r \subseteq m_\rho(r)$
- $\pi_{R_i}(m_\rho(r)) = \pi_{R_i}(r)$
- $m_\rho(m_\rho(r)) = m_\rho(r)$



- given a relation scheme **R** , a set of functional dependencies **F** , and a decomposition ρ
- how do we check that the given decomposition has a lossless join?
- there exist an algorithm that checks this in polynomial time
- **we need F as we are implicitly checking legal instances of R**

Algorithm



Algorithm - lossless join

Input: a relation scheme R , a set F of functional dependencies on R , a decomposition

$\rho = R_1, R_2, \dots, R_k$

Output: whether or not ρ has a lossless join

begin

construct a table r as follows:

r has R columns and k rows

at the intersection of the i -th row and j -th column put:

the symbol a_j if the attribute $A \in R_{i,j}$

the symbol b_{ij} otherwise

repeat

for every $X \rightarrow Y \in F$

if there are two tuples t_1 and t_2 in r such that $t_1[X] = t_2[X]$ and $t_1[Y] \neq t_2[Y]$

then for every attribute A_j in Y : if $t_1[A_j] = a_j$ then $t_2[A_j] = t_1[A_j]$

else $t_1[A_j] = t_2[A_j]$

until r has a line with all "a" or r has not changed;

if r has a row with all "a" **then:** it has a lossless join

else: it does not have a lossless join

end

a column for each attribute of R and a row for each element of the decomposition (subschema)

the attribute A_j is part of the subschema R_i

Index i = element of the decomposition = line
Index j = attribute = column

we also correctly handle the case where $t_2[A_j] = 'a_j'$

even in this case the algorithm always ends!
then we need to check if in r there is the tuple we are looking for

- we can consider the \mathbf{a}_j as particular values belonging to the domain of the attribute \mathbf{A}_j
- we can consider \mathbf{b}_{ij} as particular values belonging to the domain of the attribute \mathbf{A}_j
- we can consider **all values** \mathbf{a}_j equal to each other
- the value \mathbf{b}_{ij} is different from \mathbf{a}_j and from another value \mathbf{b}_{kj} , even if they all belong to the same domain (that of attribute \mathbf{A}_j)
- as a **consequence**, the **initial** r is a particular **instance of** the schema R



- the algorithm modifies r , so that **all** dependencies in F are **fulfilled**
- whenever it finds two tuples that are equal on the determinant but different on the dependent, it modifies the dependent so they become equal
- in doing so, it prioritizes the "a" symbol ("a" never becomes "b"; "b" can become "a")

- if two tuples have the same value in the determinant but **different values in the dependent of a** dependency, and only **one of the tuples** has the value "a" in the dependent, we change the "b" of the other tuple into an "a"
- if two tuples have the same value in the determinant but **different values in the dependent of a** dependency, and **neither** has a value of "a" in the dependent, we change one of the tuples so that they have the same value of "b" (e.g., if we have b_{ij} and b_{kj} **then we make both values b_{ij} or b_{kj}**)
- **two values are equal if they are both "a" or if they have a "b" with the same subscript**
- the algorithm stops when **ALL** pairs of tuples satisfy all the dependencies in F
- so, in the **end, r became a legal instance of R**

Theorem: let R be a relation scheme, F a set of functional dependencies on R and let $\rho = R_1, R_2, \dots, R_k$ be a decomposition of R ; the algorithm correctly decides whether ρ has a lossless join

- **Demonstration:** it must be shown that:
 - has a lossless join ($m_\rho(r)=r$, for each legal instance r)
 - if and only if
- when the algorithm terminates, r has a tuple with all "a"

- **if (by contradiction)**
 - suppose that ρ has a lossless join ($m_\rho(r) = r$) and that when the algorithm terminates the table r does not have any tuple with all "a"
 - r is a **legal instance of R** , since the algorithm terminates when there are **no more dependency violations in F**
 - as no "a" symbol appearing in the **initial value of r was changed into a "b"** by the algorithm, for each $i, i=1, \dots, k$, $\pi_{R_i}(r)$ contains (from the beginning!) a tuple with all 'a', which was the one obtained **by projecting the instance r on the attributes of R_i** , more precisely in the row corresponding to **the subschema R_i**
 - so, $m_\rho(r)$ contains a tuple with all 'a' and, consequently, $m_\rho(r) \neq r$ (**contradiction**)